In [54]:
```
# This Python 3 environment comes with many helpful analytics libraries instal
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/d
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will lis

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that ge
# You can also write temporary files to /kaggle/temp/, but they won't be saved

# Use the kagglehub client library to attach Kaggle resources like competition
# Learn more about kagglehub: https://github.com/Kaggle/kagglehub/blob/main/RE

import kagglehub
# kagglehub.dataset_download('<owner>/<dataset-slug>')
```

```
/kaggle/input/competitions/heavy-equipment-selling-price-prediction-challenge/s
ample_submission.csv
/kaggle/input/competitions/heavy-equipment-selling-price-prediction-challenge/t
rain.csv
/kaggle/input/competitions/heavy-equipment-selling-price-prediction-challenge/m
etadata.csv
/kaggle/input/competitions/heavy-equipment-selling-price-prediction-challenge/t
est.csv
```

In [55]:
```
import subprocess, sys
try:
    import catboost
except ImportError:
    subprocess.run([sys.executable, "-m", "pip", "install", "catboost", "-q"],

import numpy as np
import pandas as pd
import warnings
import os
import matplotlib.pyplot as plt
import matplotlib.ticker as mticker

warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', 60)
pd.set_option('display.float_format', '{:.2f}'.format)

from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error, r2_score
import lightgbm as lgb
from catboost import CatBoostRegressor
```

```
SEED = 42
np.random.seed(SEED)
print("ALL LIBRARIES LOADED OK.")
print("NOTE: XGBoost removed in v15 (was getting 0 percent blend weight).")
```

```
ALL LIBRARIES LOADED OK.
NOTE: XGBoost removed in v15 (was getting 0 percent blend weight).
```

In [56]:
```
def rmsle(y_true, y_pred):
    y_pred = np.maximum(y_pred, 0)
    return np.sqrt(mean_squared_error(np.log1p(y_true), np.log1p(y_pred)))

def make_submission(preds, ids, filename='submission.csv'):
    preds = np.maximum(preds, 0)
    sub = pd.DataFrame({'TransactionID': ids, 'TargetValue': preds})
    sub.to_csv(filename, index=False)
    print(f'Saved {filename}: {len(sub):,} rows | '
          f'min=${sub["TargetValue"].min():,.0f} '
          f'median=${sub["TargetValue"].median():,.0f} '
          f'max=${sub["TargetValue"].max():,.0f}')
    return sub

print("Helper functions ready.")
```

```
Helper functions ready.
```

In [57]:
```
train_df = pd.read_csv('/kaggle/input/competitions/heavy-equipment-selling-pri
test_df  = pd.read_csv('/kaggle/input/competitions/heavy-equipment-selling-pri
sample   = pd.read_csv('/kaggle/input/competitions/heavy-equipment-selling-pri

print("=" * 55)
print(f"TRAIN : {train_df.shape[0]:>8,} rows x {train_df.shape[1]} cols")
print(f"TEST  : {test_df.shape[0]:>8,} rows x {test_df.shape[1]} cols")
print("=" * 55)
t = train_df['TargetValue']
print(f"Price range: ${t.min():,.0f} - ${t.max():,.0f}")
```

```
=======================================================
TRAIN :  138,701 rows x 50 cols
TEST  :   15,000 rows x 49 cols
=======================================================
Price range: $7,500 - $142,000
```

In [58]:
```
print("Building leak-free asset history features...")
full_df = pd.concat([train_df, test_df], axis=0, ignore_index=True)
full_df = full_df.sort_values('TransactionID').reset_index(drop=True)

full_df['log_target'] = np.log1p(full_df['TargetValue'])
grouped = full_df.groupby('AssetID')

full_df['Asset_Prev_Sales_Count'] = grouped.cumcount().astype(np.int32)
full_df['Has_Asset_History']      = (full_df['Asset_Prev_Sales_Count'] > 0).as

full_df['cum_sum_log_target']      = grouped['log_target'].cumsum()
```

```
full_df['cum_sum_log_target_prev'] = full_df['cum_sum_log_target'] - full_df['
full_df['has_target']              = full_df['TargetValue'].notna().astype(np.
full_df['cum_count_target']        = grouped['has_target'].cumsum()
full_df['cum_count_target_prev']   = full_df['cum_count_target'] - full_df['ha

full_df['Asset_Prev_Mean_Price'] = (
    full_df['cum_sum_log_target_prev'] /
    full_df['cum_count_target_prev'].replace(0, np.nan)
).astype(np.float32)
full_df['Asset_Prev_Last_Price'] = grouped['log_target'].shift(1).astype(np.fl

full_df['TransactionDate_dt'] = pd.to_datetime(full_df['TransactionDate'], err
full_df['Asset_Days_Since_Last_Sale'] = (
    grouped['TransactionDate_dt'].diff().dt.days.astype(np.float32)
)

full_df = full_df.drop(columns=[
    'log_target','cum_sum_log_target','cum_sum_log_target_prev',
    'has_target','cum_count_target','cum_count_target_prev','TransactionDate_d
])

train_features = full_df[full_df['TargetValue'].notna()].copy()
test_features  = full_df[full_df['TargetValue'].isna()].copy()
train_features = (train_features.set_index('TransactionID').loc[train_df['Tran
test_features  = (test_features.set_index('TransactionID').loc[test_df['Transa

n_hist = train_features['Has_Asset_History'].sum()
print(f"Asset history features added. Train: {train_features.shape} | Test: {t
print(f"Rows with prior history: {n_hist:,} / {len(train_features):,} ({n_hist
```

```
Building leak-free asset history features...
Asset history features added. Train: (138701, 55) | Test: (15000, 55)
Rows with prior history: 19,969 / 138,701 (14.4%)
```

In [59]:

```
def preprocess_df(df):
    """
    v18 base preprocessing.
    Key v18 addition: AgeBucket categorical (5 bins) for dense AgeBucket x Sal
    All other features inherited from v16/v17.
    """
    df = df.copy()
    df['ProductConfigID'] = df['ProductConfigID'].astype(str)

    df['TransactionDate']   = pd.to_datetime(df['TransactionDate'], errors='co
    df['SaleYear']          = df['TransactionDate'].dt.year
    df['SaleMonth']         = df['TransactionDate'].dt.month
    df['SaleDay']           = df['TransactionDate'].dt.day
    df['SaleDayOfWeek']     = df['TransactionDate'].dt.dayofweek
    df['SaleDayOfYear']     = df['TransactionDate'].dt.dayofyear
    df['SaleQuarter']       = df['TransactionDate'].dt.quarter
    df['SaleWeek']          = df['TransactionDate'].dt.isocalendar().week.asty

    # Cyclical encodings (v16)
    df['SaleDayOfYear_sin'] = np.sin(2*np.pi*df['SaleDayOfYear']/365).astype(n
```

```
    df['SaleDayOfYear_cos'] = np.cos(2*np.pi*df['SaleDayOfYear']/365).astype(n
    df['SaleMonth_sin']     = np.sin(2*np.pi*df['SaleMonth']/12).astype(np.flo
    df['SaleMonth_cos']     = np.cos(2*np.pi*df['SaleMonth']/12).astype(np.flo
    df['SaleDayOfWeek_sin'] = np.sin(2*np.pi*df['SaleDayOfWeek']/7).astype(np.
    df['SaleDayOfWeek_cos'] = np.cos(2*np.pi*df['SaleDayOfWeek']/7).astype(np.
    # SaleDay cyclical (v17)
    df['SaleDay_sin']       = np.sin(2*np.pi*df['SaleDay']/31).astype(np.float
    df['SaleDay_cos']       = np.cos(2*np.pi*df['SaleDay']/31).astype(np.float

    # Time flags
    df['IsMonthEnd']        = df['TransactionDate'].dt.is_month_end.astype(np.
    df['IsMonthStart']      = df['TransactionDate'].dt.is_month_start.astype(n
    df['IsQuarterEnd']      = df['TransactionDate'].dt.is_quarter_end.astype(n
    df['IsQuarterStart']    = df['TransactionDate'].dt.is_quarter_start.astype
    df['IsYearEnd']         = df['TransactionDate'].dt.is_year_end.astype(np.i
    df['IsYearStart']       = df['TransactionDate'].dt.is_year_start.astype(np
    df = df.drop(columns=['TransactionDate'])

    # Market time trend (v17)
    df['Years_Since_2000']  = (df['SaleYear'] - 2000).clip(lower=-10, upper=20

    df.loc[df['ManufactureYear'] <= 1920, 'ManufactureYear'] = np.nan
    df.loc[df['OperationalHoursMeter'] == 0, 'OperationalHoursMeter'] = np.nan

    df['MachineAge']    = (df['SaleYear'] - df['ManufactureYear']).clip(lower=
    df['MachineAge_sq'] = (df['MachineAge'] ** 2).astype(np.float32)
    df['LogAge']        = np.log1p(df['MachineAge']).astype(np.float32)

    # NEW v18: AgeBucket categorical — 5 bins, 120 combos with SaleYear -> DEN
    # Dense TEs can be estimated accurately; sparse TEs collapse to global mea
    age_bins   = [-1, 2, 5, 10, 20, 200]
    age_labels = ['New', 'Young', 'Mid', 'Old', 'Ancient']
    df['AgeBucket'] = pd.cut(df['MachineAge'], bins=age_bins, labels=age_label
    df['AgeBucket'] = df['AgeBucket'].astype(str).fillna('Unknown').astype('ca

    tier_map  = {'Low': 0, 'Medium': 1, 'High': 2}
    scale_map = {'Mini': 0, 'Compact': 1, 'Small': 2, 'Medium': 3, 'Large / Me
    df['UtilizationTierOrd'] = df['UtilizationTier'].map(tier_map)
    df['AssetScaleOrd']      = df['AssetScaleFactor'].map(scale_map)

    hours_99 = df['OperationalHoursMeter'].quantile(0.99)
    df['OperationalHoursMeter'] = df['OperationalHoursMeter'].clip(upper=hours
    df['HoursPresent']          = df['OperationalHoursMeter'].notna().astype(n
    df['LogHours']              = np.log1p(df['OperationalHoursMeter'].fillna(
    df['UsagePerYear']          = df['OperationalHoursMeter'] / (df['MachineAg
    # LogHours x LogAge interaction (v17)
    df['LogHours_x_LogAge']     = (df['LogHours'] * df['LogAge'].fillna(0)).as

    df['IsNewMachine'] = (df['MachineAge'] <= 2).astype(np.int8)
    df['IsOldMachine'] = (df['MachineAge'] >= 20).astype(np.int8)
    df['AgeDecade']    = (df['MachineAge'] // 10).fillna(-1).astype(int)

    text_cols = df.select_dtypes(include='object').columns.tolist()
```

```
    df[text_cols] = df[text_cols].replace('None or Unspecified', np.nan)
    for col in text_cols:
        df[col] = df[col].fillna('Missing').astype('category')

    for c in ['col18', 'col19', 'UtilizationTier', 'AssetScaleFactor']:
        if c in df.columns:
            df = df.drop(columns=[c])

    return df

print("Preprocessing features...")
train_proc = preprocess_df(train_features)
test_proc  = preprocess_df(test_features)
print(f"Train: {train_proc.shape} | Test: {test_proc.shape}")
```

```
Preprocessing features...
Train: (138701, 85) | Test: (15000, 85)
```

In [60]:
```
ID_COLS = ['TransactionID', 'AssetID']

y_all        = np.log1p(train_proc['TargetValue'].values)
X_all        = train_proc.drop(columns=ID_COLS + ['TargetValue'], errors='ignc
X_test_final = test_proc.drop(columns=ID_COLS + ['TargetValue'], errors='ignor
test_ids     = test_df['TransactionID'].values
cat_cols     = X_all.select_dtypes(include=['category', 'object']).columns.tol

print(f"Feature matrix : {X_all.shape}")
print(f"Test matrix    : {X_test_final.shape}")
print(f"Cat features   : {len(cat_cols)}")
```

```
Feature matrix : (138701, 82)
Test matrix    : (15000, 82)
Cat features   : 41
```

In [61]:
```
def add_oof_target_encoding(X_tr, y_tr, X_vl, X_te, cols, n_splits=5, alpha=15
    X_tr, X_vl, X_te = X_tr.copy(), X_vl.copy(), X_te.copy()
    global_mean = float(y_tr.mean())
    kf_inner    = KFold(n_splits=n_splits, shuffle=True, random_state=seed)
    for col in cols:
        col_te = col + '_TE'
        X_tr[col_te] = X_vl[col_te] = X_te[col_te] = global_mean
        tr_val = X_tr[col].astype(str)
        vl_val = X_vl[col].astype(str)
        te_val = X_te[col].astype(str)
        stats = pd.DataFrame({'cat': tr_val, 'y': y_tr})
        agg   = stats.groupby('cat')['y'].agg(['mean', 'count'])
        agg['sm'] = (agg['count']*agg['mean'] + alpha*global_mean) / (agg['cou
        m = agg['sm'].to_dict()
        X_vl[col_te] = vl_val.map(m).fillna(global_mean).values
        X_te[col_te] = te_val.map(m).fillna(global_mean).values
        col_idx = X_tr.columns.get_loc(col_te)
        for tr_in, vl_in in kf_inner.split(X_tr):
            in_s = pd.DataFrame({'cat': tr_val.iloc[tr_in], 'y': y_tr[tr_in]})
            in_a = in_s.groupby('cat')['y'].agg(['mean', 'count'])
```

```
            in_a['sm'] = (in_a['count']*in_a['mean'] + alpha*global_mean) / (i
            X_tr.iloc[vl_in, col_idx] = tr_val.iloc[vl_in].map(in_a['sm'].to_d
    return X_tr, X_vl, X_te


def add_interaction_te(X_tr, y_tr, X_vl, X_te, col_a, col_b, alpha=15.0, seed=
    """
    Interaction Target Encoding with OOF inner-fold smoothing.
    Alpha should be set proportional to sparsity:
      - Dense (avg>100/cell): alpha=10-15
      - Moderate (avg 30-100/cell): alpha=15-25
      - Sparse (<30/cell): DO NOT USE (noise only)
    """
    X_tr, X_vl, X_te = X_tr.copy(), X_vl.copy(), X_te.copy()
    global_mean = float(y_tr.mean())
    kf_inner    = KFold(n_splits=5, shuffle=True, random_state=seed)
    feat_name   = col_a + '__x__' + col_b + '_TE'
    tr_int = X_tr[col_a].astype(str) + '__' + X_tr[col_b].astype(str)
    vl_int = X_vl[col_a].astype(str) + '__' + X_vl[col_b].astype(str)
    te_int = X_te[col_a].astype(str) + '__' + X_te[col_b].astype(str)
    X_tr[feat_name] = X_vl[feat_name] = X_te[feat_name] = global_mean
    stats = pd.DataFrame({'cat': tr_int, 'y': y_tr})
    agg   = stats.groupby('cat')['y'].agg(['mean', 'count'])
    agg['sm'] = (agg['count']*agg['mean'] + alpha*global_mean) / (agg['count']
    m = agg['sm'].to_dict()
    X_vl[feat_name] = vl_int.map(m).fillna(global_mean).values
    X_te[feat_name] = te_int.map(m).fillna(global_mean).values
    col_idx = X_tr.columns.get_loc(feat_name)
    for tr_in, vl_in in kf_inner.split(X_tr):
        in_s = pd.DataFrame({'cat': tr_int.iloc[tr_in], 'y': y_tr[tr_in]})
        in_a = in_s.groupby('cat')['y'].agg(['mean', 'count'])
        in_a['sm'] = (in_a['count']*in_a['mean'] + alpha*global_mean) / (in_a[
        X_tr.iloc[vl_in, col_idx] = tr_int.iloc[vl_in].map(in_a['sm'].to_dict(
    return X_tr, X_vl, X_te


def add_fold_group_features(X_tr, X_vl, X_te):
    X_tr, X_vl, X_te = X_tr.copy(), X_vl.copy(), X_te.copy()
    freq_cols = ['ProductConfigID', 'Spec_FullDescriptor', 'Spec_BaseClass', '
    for col in freq_cols:
        freq_map = X_tr[col].astype(str).value_counts().to_dict()
        for df_part in [X_tr, X_vl, X_te]:
            df_part[col+'_Freq'] = df_part[col].astype(str).map(freq_map).fill
    group_cols   = ['InventoryGroupCategory', 'Spec_BaseClass', 'ProductConfig
                    'VendorPartnerID', 'RegionCode', 'FunctionalClassificatior
    target_feats = ['OperationalHoursMeter', 'MachineAge']
    for grp_col in group_cols:
        for target_col in target_feats:
            global_mean = float(X_tr[target_col].mean())
            group_means = X_tr.groupby(X_tr[grp_col].astype(str))[target_col].
            feat_name   = f'{grp_col}_{target_col}_mean'
            for df_part in [X_tr, X_vl, X_te]:
                df_part[feat_name] = df_part[grp_col].astype(str).map(group_me
```

```
            for df_part in [X_tr, X_vl, X_te]:
                df_part[f'{target_col}_rel_{grp_col}'] = (df_part[target_col]/
    return X_tr, X_vl, X_te


TE_COLS = [
    'Spec_FullDescriptor', 'Spec_BaseClass', 'ProductConfigID',
    'FunctionalClassification', 'Spec_SubClass', 'Spec_VariantModifier', 'Regi
]

print("Feature engineering functions compiled.")
print("Density-checked interaction TEs for v18:")
print("  SAFE (avg>100/cell): AgeBucket x SaleYear, InventoryGroupCategory x S
print("  OK   (avg~112/cell): Spec_SubClass x SaleYear")
```

```
Feature engineering functions compiled.
Density-checked interaction TEs for v18:
  SAFE (avg>100/cell): AgeBucket x SaleYear, InventoryGroupCategory x SaleYear,
RegionCode x SaleYear
  OK   (avg~112/cell): Spec_SubClass x SaleYear
```

In [62]:

```
N_FOLDS = 5
kf = KFold(n_splits=N_FOLDS, shuffle=True, random_state=SEED)

oof_lgb1 = np.zeros(len(X_all))
oof_lgb2 = np.zeros(len(X_all))
oof_cat  = np.zeros(len(X_all))

test_lgb1_folds = np.zeros((len(X_test_final), N_FOLDS))
test_lgb2_folds = np.zeros((len(X_test_final), N_FOLDS))
test_cat_folds  = np.zeros((len(X_test_final), N_FOLDS))

fold_scores = []

print("Starting 5-Fold CV - v18 (Dense Interactions + More Capacity)")
print("=" * 70)

for fold, (tr_idx, vl_idx) in enumerate(kf.split(X_all), 1):
    X_tr_raw = X_all.iloc[tr_idx].copy()
    X_vl_raw = X_all.iloc[vl_idx].copy()
    y_tr     = y_all[tr_idx]
    y_vl     = y_all[vl_idx]

    # Step 1: 7 single-column OOF TEs (same since v15)
    X_tr, X_vl, X_te = add_oof_target_encoding(
        X_tr_raw, y_tr, X_vl_raw, X_test_final.copy(), TE_COLS, alpha=15.0
    )

    # Step 2: Region interaction TEs (3 from v15/v16, proven)
    for ca, cb, alp in [
        ('InventoryGroupDescription', 'RegionCode', 20.0),
        ('FunctionalClassification',  'RegionCode', 22.0),
        ('Spec_SubClass',             'RegionCode', 22.0),
    ]:
```

```
        X_tr, X_vl, X_te = add_interaction_te(X_tr, y_tr, X_vl, X_te, ca, cb,

    # Step 3: Time interaction TE - only FunctionalClass x SaleYear kept from
    # (it was the only dense one: 102 samples/cell. Others had 6.8 and 14.6 -
    X_tr, X_vl, X_te = add_interaction_te(
        X_tr, y_tr, X_vl, X_te,
        'FunctionalClassification', 'SaleYear', alpha=10.0
    )

    # Step 4: NEW v18 DENSE interactions (all verified >= 110 samples/cell)
    # 4a: AgeBucket x SaleYear (924/cell) - how does each age group depreciate
    X_tr, X_vl, X_te = add_interaction_te(
        X_tr, y_tr, X_vl, X_te,
        'AgeBucket', 'SaleYear', alpha=10.0
    )
    # 4b: InventoryGroupCategory x SaleYear (924/cell) - how does each equipme
    #     change price year-over-year?
    X_tr, X_vl, X_te = add_interaction_te(
        X_tr, y_tr, X_vl, X_te,
        'InventoryGroupCategory', 'SaleYear', alpha=10.0
    )
    # 4c: RegionCode x SaleYear (110/cell) - regional price trends over time
    X_tr, X_vl, X_te = add_interaction_te(
        X_tr, y_tr, X_vl, X_te,
        'RegionCode', 'SaleYear', alpha=12.0
    )
    # 4d: Spec_SubClass x SaleYear (112/cell) - equipment sub-class price tren
    X_tr, X_vl, X_te = add_interaction_te(
        X_tr, y_tr, X_vl, X_te,
        'Spec_SubClass', 'SaleYear', alpha=15.0
    )

    # Step 5: Group/frequency features
    X_tr, X_vl, X_te = add_fold_group_features(X_tr, X_vl, X_te)

    cat_fold = [c for c in cat_cols if c in X_tr.columns]

    def to_str(df, cols):
        df = df.copy()
        for c in cols:
            if c in df.columns:
                df[c] = df[c].astype(str)
        return df

    X_tr_c = to_str(X_tr, cat_fold)
    X_vl_c = to_str(X_vl, cat_fold)
    X_te_c = to_str(X_te, cat_fold)

    # LightGBM 1 — num_leaves=1023 (increased from 511 for more capacity)
    lgb1_m = lgb.LGBMRegressor(
        n_estimators=15000, max_depth=-1, num_leaves=1023,
        learning_rate=0.010, subsample=0.75, colsample_bytree=0.70,
        min_child_samples=30, reg_alpha=0.1, reg_lambda=3.0,
```

```
        random_state=42, n_jobs=-1, verbose=-1
    )
    lgb1_m.fit(X_tr, y_tr, eval_set=[(X_vl, y_vl)],
               callbacks=[lgb.early_stopping(300, verbose=False), lgb.log_eval
    oof_lgb1[vl_idx]           = lgb1_m.predict(X_vl)
    test_lgb1_folds[:, fold-1] = lgb1_m.predict(X_te)

    # LightGBM 2 — num_leaves=1023, diverse seed/colsample
    lgb2_m = lgb.LGBMRegressor(
        n_estimators=15000, max_depth=-1, num_leaves=1023,
        learning_rate=0.010, subsample=0.80, colsample_bytree=0.65,
        min_child_samples=30, reg_alpha=0.2, reg_lambda=2.0,
        random_state=123, n_jobs=-1, verbose=-1
    )
    lgb2_m.fit(X_tr, y_tr, eval_set=[(X_vl, y_vl)],
               callbacks=[lgb.early_stopping(300, verbose=False), lgb.log_eval
    oof_lgb2[vl_idx]           = lgb2_m.predict(X_vl)
    test_lgb2_folds[:, fold-1] = lgb2_m.predict(X_te)

    # CatBoost — max_leaves=255 (kept, CatBoost handles categoricals natively
    cat_m = CatBoostRegressor(
        iterations=15000, learning_rate=0.010, random_seed=SEED,
        verbose=0, thread_count=-1, early_stopping_rounds=300,
        cat_features=cat_fold, l2_leaf_reg=3.0, bagging_temperature=0.5,
        grow_policy='Lossguide', max_leaves=255, min_data_in_leaf=30
    )
    cat_m.fit(X_tr_c, y_tr, eval_set=(X_vl_c, y_vl))
    oof_cat[vl_idx]           = cat_m.predict(X_vl_c)
    test_cat_folds[:, fold-1] = cat_m.predict(X_te_c)

    s_1 = rmsle(np.expm1(y_vl), np.expm1(oof_lgb1[vl_idx]))
    s_2 = rmsle(np.expm1(y_vl), np.expm1(oof_lgb2[vl_idx]))
    s_c = rmsle(np.expm1(y_vl), np.expm1(oof_cat[vl_idx]))
    fold_scores.append((s_1, s_2, s_c))
    print(f"Fold {fold}: LGB1={s_1:.5f}(iter={lgb1_m.best_iteration_:4d}) | "
          f"LGB2={s_2:.5f}(iter={lgb2_m.best_iteration_:4d}) | CAT={s_c:.5f}")

oof_1 = rmsle(np.expm1(y_all), np.expm1(oof_lgb1))
oof_2 = rmsle(np.expm1(y_all), np.expm1(oof_lgb2))
oof_c = rmsle(np.expm1(y_all), np.expm1(oof_cat))
print("=" * 70)
print(f"LGB1 OOF RMSLE : {oof_1:.5f}  (v17 was 0.19542, need improvement)")
print(f"LGB2 OOF RMSLE : {oof_2:.5f}  (v17 was 0.19527)")
print(f"CAT  OOF RMSLE : {oof_c:.5f}  (v17 was 0.20332)")
```

```
Starting 5-Fold CV - v18 (Dense Interactions + More Capacity)
======================================================================
Fold 1: LGB1=0.19265(iter=3992) | LGB2=0.19234(iter=4107) | CAT=0.20089
Fold 2: LGB1=0.19562(iter=3815) | LGB2=0.19534(iter=4145) | CAT=0.20391
Fold 3: LGB1=0.19594(iter=2950) | LGB2=0.19579(iter=3485) | CAT=0.20455
Fold 4: LGB1=0.19630(iter=3410) | LGB2=0.19616(iter=3324) | CAT=0.20458
Fold 5: LGB1=0.19492(iter=2748) | LGB2=0.19488(iter=3121) | CAT=0.20219
======================================================================
LGB1 OOF RMSLE : 0.19509  (v17 was 0.19542, need improvement)
LGB2 OOF RMSLE : 0.19491  (v17 was 0.19527)
CAT  OOF RMSLE : 0.20323  (v17 was 0.20332)
```

In [63]:
```
print("Searching optimal blend weights...")
best_blend, best_w = 999.0, (0.30, 0.40, 0.30)

for w1 in np.arange(0.0, 1.01, 0.05):
    for w2 in np.arange(0.0, 1.01 - w1, 0.05):
        w3    = max(round(1.0 - w1 - w2, 6), 0.0)
        score = rmsle(np.expm1(y_all),
                      np.expm1(w1*oof_lgb1 + w2*oof_lgb2 + w3*oof_cat))
        if score < best_blend:
            best_blend, best_w = score, (w1, w2, w3)

print("=" * 70)
print(f"Best blend    : LGB1={best_w[0]:.2f}  LGB2={best_w[1]:.2f}  CAT={best_
print(f"Blended RMSLE : {best_blend:.5f}  (v17 blended was 0.19359)")
target = 0.18757
print(f"Kaggle v17 score : {target}  ->  {'[BEAT v17!]' if best_blend < 0.1935
```

```
Searching optimal blend weights...
======================================================================
Best blend    : LGB1=0.25  LGB2=0.45  CAT=0.30
Blended RMSLE : 0.19326  (v17 blended was 0.19359)
Kaggle v17 score : 0.18757  ->  [BEAT v17!]
```

In [64]:
```
w1, w2, w3 = best_w
test_blend   = (w1*test_lgb1_folds.mean(1) + w2*test_lgb2_folds.mean(1) + w3*t
final_prices = np.expm1(test_blend)
plo = train_df['TargetValue'].quantile(0.001)
phi = train_df['TargetValue'].quantile(0.999)
final_prices = np.clip(final_prices, plo, phi)
print(f"Price clipped to [${plo:,.0f}, ${phi:,.0f}]")
sub = make_submission(final_prices, test_ids, 'submission.csv')
print()
print(sub.head())
```

```
Price clipped to [$7,500, $137,500]
Saved submission.csv: 15,000 rows | min=$7,728 median=$34,615 max=$137,500

   TransactionID  TargetValue
0        1139307     69296.52
1        1139419     81322.36
2        1139482     30853.24
3        1139522     20357.37
4        1139684     12913.14
```

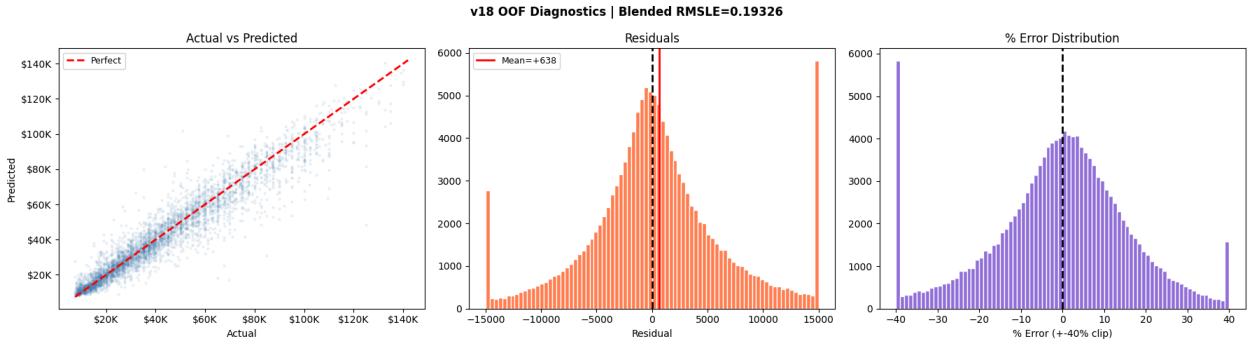
In [65]:
```
oof_blend  = best_w[0]*oof_lgb1 + best_w[1]*oof_lgb2 + best_w[2]*oof_cat
pred_price = np.expm1(oof_blend)
true_price = np.expm1(y_all)
residuals  = true_price - pred_price
pct_err    = (residuals / true_price) * 100

fig, axes = plt.subplots(1, 3, figsize=(18, 5))
fig.suptitle(f'v18 OOF Diagnostics | Blended RMSLE={best_blend:.5f}', fontsize

idx = np.random.choice(len(true_price), min(5000, len(true_price)), replace=Fa
axes[0].scatter(true_price[idx], pred_price[idx], alpha=0.07, s=4, color='stee
mn, mx = true_price.min(), true_price.max()
axes[0].plot([mn, mx], [mn, mx], 'r--', lw=2, label='Perfect')
axes[0].set_xlabel('Actual'); axes[0].set_ylabel('Predicted'); axes[0].set_tit
axes[0].legend(fontsize=9)
fmt = lambda x, _: f'${x/1000:.0f}K'
axes[0].xaxis.set_major_formatter(mticker.FuncFormatter(fmt))
axes[0].yaxis.set_major_formatter(mticker.FuncFormatter(fmt))
axes[1].hist(residuals.clip(-15000, 15000), bins=80, color='coral', edgecolor=
axes[1].axvline(0, color='black', lw=2, ls='--')
axes[1].axvline(residuals.mean(), color='red', lw=2, label=f'Mean={residuals.m
axes[1].set_xlabel('Residual'); axes[1].set_title('Residuals'); axes[1].legend
axes[2].hist(pct_err.clip(-40, 40), bins=80, color='mediumpurple', edgecolor='
axes[2].axvline(0, color='black', lw=2, ls='--')
axes[2].set_xlabel('% Error (+-40% clip)'); axes[2].set_title('% Error Distrib
plt.tight_layout()
plt.show()

print(f"Blended OOF RMSLE : {best_blend:.5f}")
print(f"R2                : {r2_score(true_price, pred_price):.4f}")
print(f"MAE               : ${np.mean(np.abs(residuals)):,.0f}")
print(f"Median abs % err  : {np.median(np.abs(pct_err)):.1f}%")

fi = pd.Series(lgb1_m.feature_importances_, index=X_tr.columns).sort_values(as
print("\nTop 40 features (LGB1 last fold):")
print(fi.head(40).to_string())
```



```
Blended OOF RMSLE : 0.19326
R2                : 0.9161
MAE               : $5,042
Median abs % err  : 9.7%

Top 40 features (LGB1 last fold):
RegionCode__x__SaleYear_TE                              116945
MachineAge_rel_ProductConfigID                           93290
InventoryGroupCategory__x__SaleYear_TE                   92179
RegionCode                                               90712
AgeBucket__x__SaleYear_TE                                87030
SaleDayOfYear_cos                                        83645
SaleDayOfYear_sin                                        82577
ProductConfigID                                          78914
SaleDayOfYear                                            78433
Spec_SubClass__x__SaleYear_TE                            76306
MachineAge_rel_Spec_BaseClass                            71187
SaleDay_sin                                              68030
FunctionalClassification__x__SaleYear_TE                 67657
MachineAge_rel_RegionCode                                66060
Spec_SubClass__x__RegionCode_TE                          63401
RegionCode_TE                                            62330
SaleDay                                                  62069
FunctionalClassification__x__RegionCode_TE               61935
MachineAge_rel_FunctionalClassification                  61094
InventoryGroupDescription__x__RegionCode_TE              60172
MachineAge_rel_VendorPartnerID                           57043
SaleDay_cos                                              53257
ManufactureYear                                          47772
Spec_BaseClass                                           44652
SaleYear                                                 41302
Spec_FullDescriptor                                      40846
Spec_FullDescriptor_TE                                   38747
Spec_VariantModifier_TE                                  37642
Spec_BaseClass_TE                                        37408
FunctionalClassification_TE                              35570
Spec_SubClass_TE                                         34213
MachineAge_rel_InventoryGroupCategory                    34030
OperationalHoursMeter_rel_ProductConfigID                32216
ProductConfigID_Freq                                     30833
FunctionalClassification                                 30721
OperationalHoursMeter_rel_Spec_BaseClass                 26066
UsagePerYear                                             25295
RegionCode_OperationalHoursMeter_mean                    24703
OperationalHoursMeter_rel_FunctionalClassification       24225
SaleDayOfWeek                                            23319
```